

Applied Computational Multi-Omics

Genetic Alterations and Functional Impact

Daniel Sobral, INSA

17/03/2025



MCBBi | **NOVA**
UNIVERSIDADE NOVA
DE LISBOA

Genetic Alterations and Functional Impact

Learning Objectives



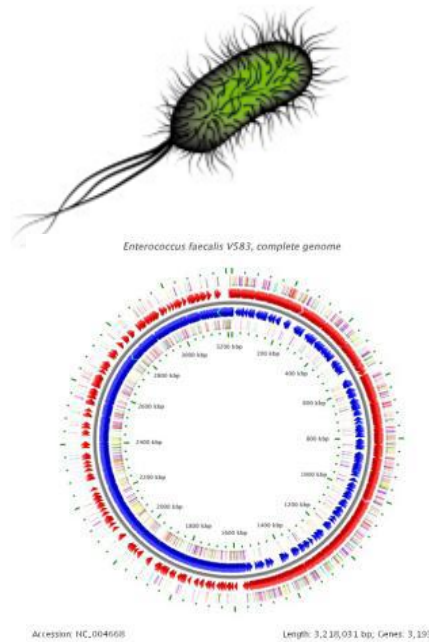
Introduction to variant calling

- 2 major classes: SNPs / indels and large structural variants
- Factors that influence the variant calling process
- Overview of the VCF file format
- Variant quality and Genotype quality

Introduction to Variant Annotation


Variant Calling

Common question: find mutations underlying phenotypes


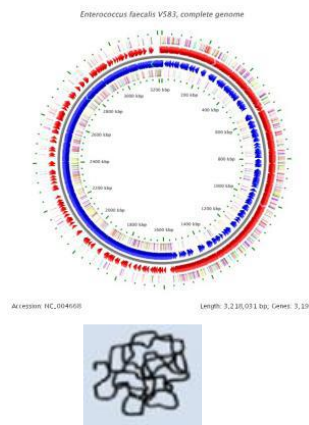


Data pre-processing for variant discovery


DNA Extraction



Whole Genome VS Targeted



DNA Extraction



Eg. TruSight One Enrichment Panel (Human)
Eg. SARS-CoV-2 ARTIC Amplicon panel

Note: Amplicon VS Enrichment

How to calculate coverage

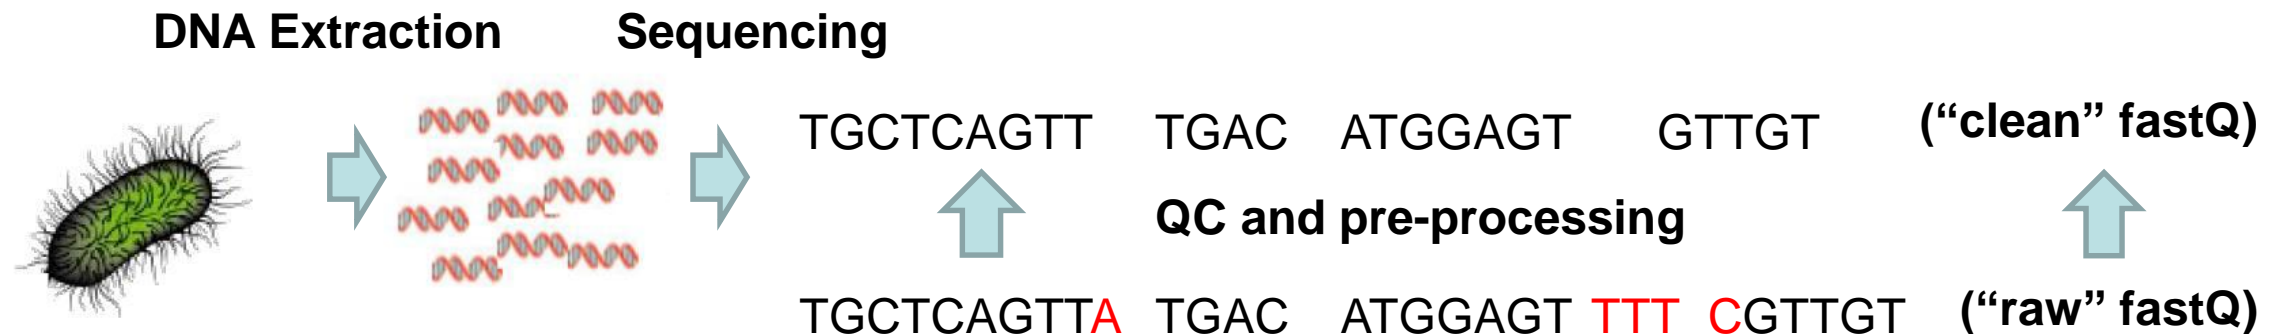
Example for the Human Genome ($\sim 3 \times 10^9$ bp):

- WGS 30x coverage, 150bp read pairs
 - $30 \times 3 \times 10^9 / (150 \times 2) = 3 \times 10^8 = 300$ million read pairs
- WES ($\sim 1\text{-}2\%$ genome) 30x coverage, 150bp read pairs
 - 6 million read pairs (theoretical minimum, but usually more)

Data pre-processing for variant discovery



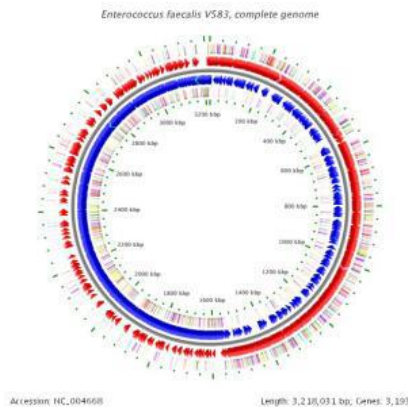
Data pre-processing for variant discovery



Data pre-processing for variant discovery

(fasta file)



Reference Genome



geneA geneB geneC
AAGCGATGACTGCATGCACAGTTGTGTGTTTTCACGTGAC

DNA Extraction

Sequencing



TGCTCAGTT

TGAC

ATGGAGT

GTTGT

("clean" fastQ)

QC and pre-processing

TGCTCAGTTA

TGAC

ATGGAGT

TTT


CGTTGT

("raw" fastQ)

Data pre-processing for variant discovery

(fasta file)



Reference Genome



geneA geneB geneC
AAGCGATGACTGCATGCACAGTTGTGTGTTTTCACGTGAC
TGCTCAGTT

DNA Extraction

Sequencing



TGCTCAGTT



TGAC

ATGGAGT

GTTGT

("clean" fastQ)

QC and pre-processing



TGCTCAGTTA

TGAC

ATGGAGT

TTT

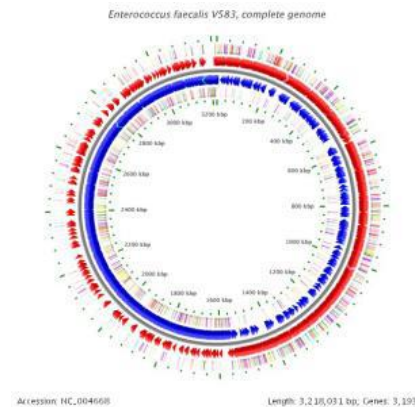
CGTTGT

("raw" fastQ)

Data pre-processing for variant discovery

(fasta file)

Reference Genome



geneA geneB geneC



AAGCGATGACTGCATGCACAGTTGTGTGTTTTCACGTGAC

TGCTCAGTT

DNA Extraction

Sequencing

Alignment



TGCTCAGTT

TGAC

ATGGAGT

GTTGT

("clean" fastQ)

QC and pre-processing



TGCTCAGTTA

TGAC

ATGGAGT

TTT


CGTTGT

("raw" fastQ)

Data pre-processing for variant discovery

(fasta file)

Reference Genome





ACTCCAT GTTGT
TGCTCAGTT

geneA geneB geneC
AAGCGATGACTGCATGCACAGTTGTGTGTTTTCACGTGAC

DNA Extraction

Sequencing

Alignment



TGCTCAGTT

TGAC

ATGGAGT

GTTGT

("clean" fastQ)

QC and pre-processing

TGCTCAGTTA

TGAC

ATGGAGT

TTT


CGTTGT

("raw" fastQ)

Data pre-processing for variant discovery

(fasta file)

Reference Genome



TCCATGC AGTTGTGT
 ACTCCAT GTTGT
 AAGCGATG TGCTCAGTT GTGTGTTT-CA
 geneA geneB geneC
 AAGCGATGACTGCATGCACAGTTGTGTGTTTTCACGTGAC

(SAM/BAM)


DNA Extraction

Sequencing

Alignment

("clean" fastQ)

("raw" fastQ)



TGCTCAGTT TGAC ATGGAGT GTTGT
 QC and pre-processing
 TGCTCAGTTA TGAC ATGGAGT TTT CGTTGT

(ACTCCAT)

SAM/BAM format

A file format to represent alignments

BAM -> binary form of SAM

```

Coor      12345678901234 5678901234567890123456789012345
ref       AGCATGTTAGATAA**GATAGCTGTGCTAGTAGGCAGTCAGCGCCAT

+r001/1      TTAGATAAAAGGATA*CTG
+r002        aaaAGATAA*GGATA
+r003        gcctaAGCTAA
+r004                ATAGCT.....TCAGC
-r003                ttagctTAGGC
-r001/2                CAGCGGCAT
    
```



```

@HD VN:1.6 SO:coordinate
@SQ SN:ref LN:45
r001 99 ref 7 30 8M2I4M1D3M = 37 39 TTAGATAAAAGGATACTG
r002 0 ref 9 30 3S6M1P1I4M * 0 0 AAAAGATAAGGATA
r003 0 ref 9 30 5S6M * 0 0 GCCTAAGCTAA
r004 0 ref 16 30 6M14N5M * 0 0 ATAGCTTCAGC
r003 2064 ref 29 17 6H5M * 0 0 TAGGC
r001 147 ref 37 30 9M = 7 -39 CAGCGGCAT
    
```

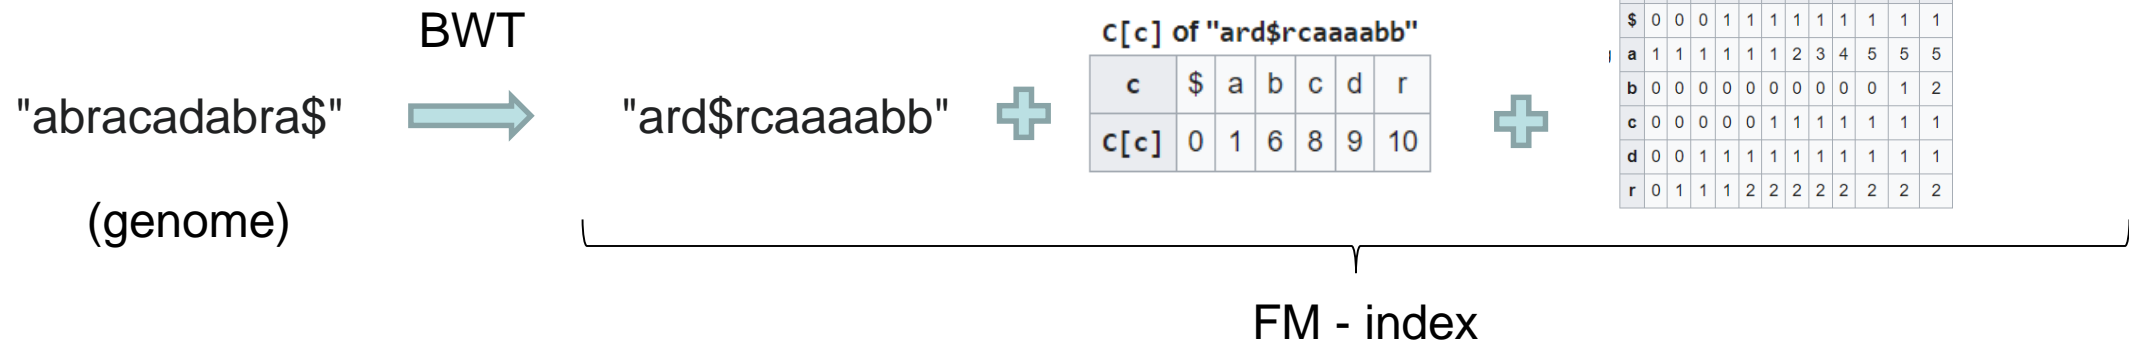
<https://samtools.github.io/hts-specs/SAMv1.pdf>

How alignment is made in practice

BWA (Burrows-Wheeler Aligner) is the most popular tool for WGS/WES

- Align millions of short reads to a human-sized genome in minutes

It is based on the FM-index of the Burrows-Wheeler Transform



<https://en.wikipedia.org/wiki/FM-index>


How alignment is made in practice

Genome BWT FM-index needs to be created (only once)

\$ bwa index genome.fasta

(genome.fasta file)

Reference Genome



BWT



c	\$	a	b	c	d	r
c[c]	0	1	6	8	9	10



	a	r	d	\$	r	c	a	a	a	a	b	b
1	2	3	4	5	6	7	8	9	10	11	12	
\$	0	0	0	1	1	1	1	1	1	1	1	1
a	1	1	1	1	1	1	2	3	4	5	5	5
b	0	0	0	0	0	0	0	0	0	1	2	
c	0	0	0	0	0	1	1	1	1	1	1	1
d	0	0	1	1	1	1	1	1	1	1	1	1
r	0	1	1	1	2	2	2	2	2	2	2	2

3×10^9 bases

AAGCGATGACTGCATGCACAGTTGTGTGTTTCACGTGAC

FM - index

How alignment is made in practice

Basic command to generate alignments with BWA:

```
$ bwa mem genome.fasta reads_R1.fastq(.gz) reads_R2.fastq(.gz)>output.sam
```

reads_R1.fastq
TGCTCAGTT

reads_R2.fastq
ACGTCCGA

BWT

c	\$	a	b	c	d	r
c[c]	0	1	6	8	9	10

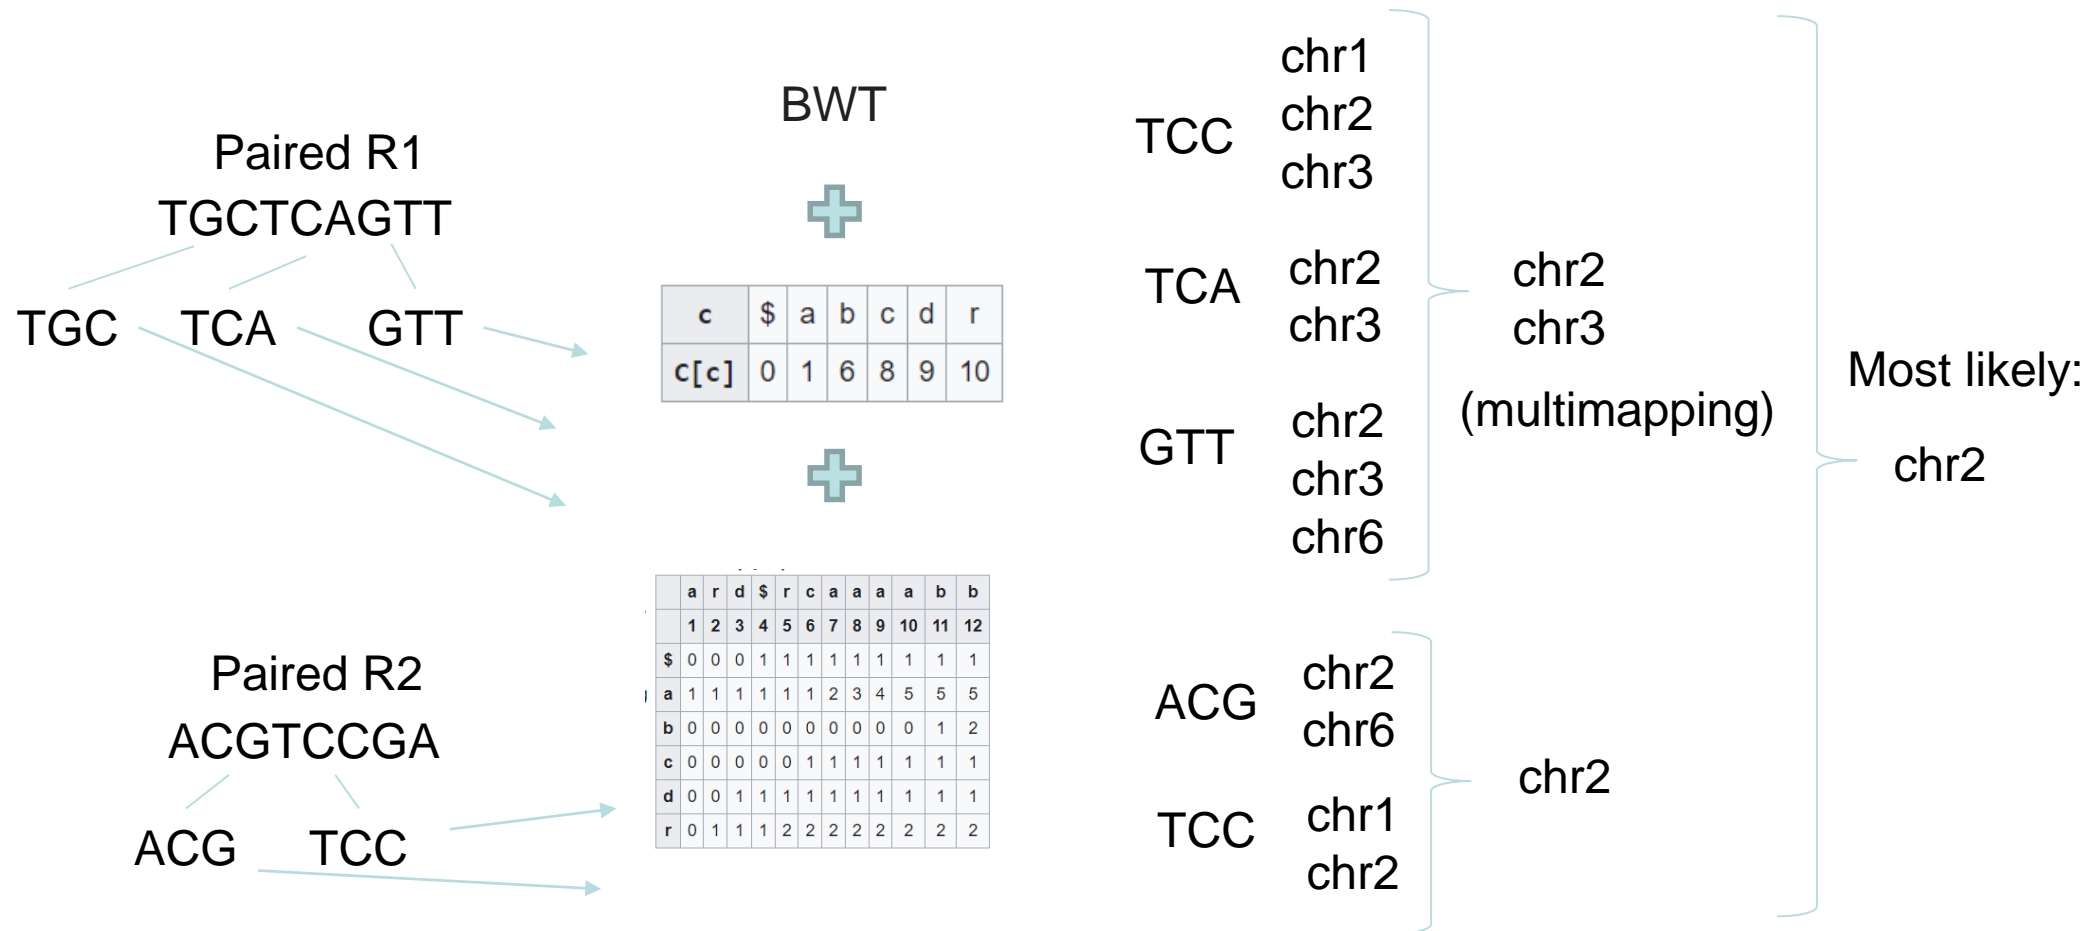
a	r	d	\$	r	c	a	a	a	a	b	b
1	2	3	4	5	6	7	8	9	10	11	12
\$	0	0	0	1	1	1	1	1	1	1	1
a	1	1	1	1	1	1	2	3	4	5	5
b	0	0	0	0	0	0	0	0	0	1	2
c	0	0	0	0	0	1	1	1	1	1	1
d	0	0	1	1	1	1	1	1	1	1	1
r	0	1	1	1	2	2	2	2	2	2	2

Output in SAM format

TGCTCAGTT Chr2 position xxx
ACGTCCGA Chr2 position yyy

```
@HD VN:1.6 SO:coordinate
@SQ SN:ref LN:45
r001 99 ref 7 30 8M2I4M1D3M = 37 39 TTAGATAAAGGATACTG
r002 0 ref 9 30 3S6M1P1I4M * 0 0 AAAAGATAAGGATA
r003 0 ref 9 30 5S6M * 0 0 GCCTAAGCTAA
r004 0 ref 16 30 6M14N5M * 0 0 ATAGCTTCAGC
r003 2064 ref 29 17 6H5M * 0 0 TAGGC
r001 147 ref 37 30 9M = 7 -39 CAGCGGCAT
```

How alignment is made in practice






How alignment is made in practice: Summary

- Special algorithms are used to have fast alignments
 - They are not guaranteed to be perfect but most of the time they are very good
- A read can map equally well to multiple regions (multimappings)
 - BWA reports **one primary alignment (randomly chosen) with mapping quality of 0**
 - Depending on the software, it can generate **secondary alignments**
 - Information of paired reads are used to disambiguate multimappings if possible
- Alignments are made piece-wise (a read is split in segments)
 - A read alignment can be split in a primary and **supplementary alignment(s)**
 - Eg. splicing in RNA-Seq; large deletions
 - Sometimes, only a part of the read is aligned (the rest is “masked”/hidden)
 - Particularly in repetitive areas this can lead to false alignments

Variant Calling

(fasta file)

Reference Genome



TCCATGC AGTTGTGT
 ACTCCAT GTTGT
 AAGCGATG TGCTCAGTT GTGTGTTT-CA
 geneA geneB geneC
 AAGCGATGACTGCATGCACAGTTGTGTGTTTTCACGTGAC

DNA Extraction


Sequencing

Alignment

(SAM/BAM)

("clean" fastQ)

("raw" fastQ)



TGCTCAGTT

TGAC

ATGGAGT

GTTGT

QC and pre-processing

TGCTCAGTTA

TGAC

ATGGAGT


TTT

CGTTGT

(ACTCCAT)

Variant Calling

(fasta file)
Reference Genome



SNV SNV


TCCATGC AGTTGTGT
ACTCCAT GTTGT

AAGCGATG TGCTCAGTT GTGTGTTT-CA


geneA geneB geneC

AAGCGATGACTGCATGCACAGTTGTGTGTTTTCACGTGAC

DNA Extraction



Sequencing



TGCTCAGTT

TGAC

ATGGAGT

GTTGT

QC and pre-processing

TGCTCAGTTA

TGAC

ATGGAGT

TTT

CGTTGT

("clean" fastQ)


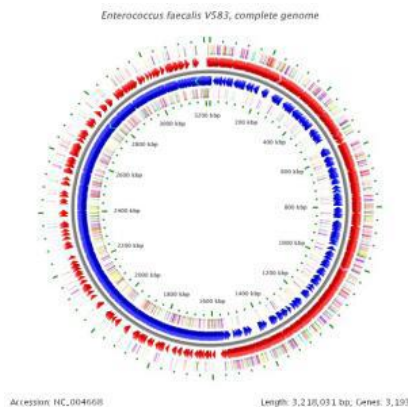
("raw" fastQ)

Alignment

(SAM/BAM)



Variant Calling

(fasta file)
Reference Genome





Variant Calling

(fasta file)
Reference Genome



Variant Calling

(fasta file)
Reference Genome



The VCF format

```
##fileformat=VCFv4.4
##fileDate=20090805
##source=myImputationProgramV3.1
##reference=file:///seq/references/1000GenomesPilot-NCBI36.fasta
##contig=<ID=20,length=62435964,assembly=B36,md5=f126cdf8a6e0c7f379d618ff66beb2da,species="Homo sapiens",taxonomy=x>
##phasing=partial
##INFO=<ID=NS,Number=1,Type=Integer,Description="Number of Samples With Data">
##INFO=<ID=DP,Number=1,Type=Integer,Description="Total Depth">
##INFO=<ID=AF,Number=A,Type=Float,Description="Allele Frequency">
##INFO=<ID=AA,Number=1,Type=String,Description="Ancestral Allele">
##INFO=<ID=DB,Number=0,Type=Flag,Description="dbSNP membership, build 129">
##INFO=<ID=H2,Number=0,Type=Flag,Description="HapMap2 membership">
##FILTER=<ID=q10,Description="Quality below 10">
##FILTER=<ID=s50,Description="Less than 50% of samples have data">
##FORMAT=<ID=GT,Number=1,Type=String,Description="Genotype Quality">
##FORMAT=<ID=DP,Number=1,Type=Integer,Description="Read Depth">
##FORMAT=<ID=HQ,Number=2,Type=Integer,Description="Haplotype Quality">
```

#CHROM	POS	ID	REF	ALT	QUAL	FILTER	INFO
20	14370	rs6054257	G	A	29	PASS	NS=3;DP=14;AF=0.5;DB;H2
20	17330	.	T	A	3	q10	NS=3;DP=11;AF=0.017
20	1110696	rs6040355	A	G,T	67	PASS	NS=2;DP=10;AF=0.333,0.667;AA=T;DB
20	1230237	.	T	.	47	PASS	NS=3;DP=13;AA=T
20	1234567	microsat1	GTC	G,GTCT	50	PASS	NS=3;DP=9;AA=G

CHROM	chromosome
POS	position of the start of the variant
ID	unique identifier of the variant (e.g. rs number for SNPs)
REF	reference allele
ALT	comma separated list of alternate non-reference alleles
QUAL	phred-scaled quality score
FILTER	site filtering information
INFO	user extensible annotation (e.g. samtools and GATK may differ in this)

FORMAT	NA00001	NA00002	NA00003
GT:GQ:DP:HQ	0 0:48:1:51,51	1 0:48:8:51,51	1/1:43:5:.,.
GT:GQ:DP:HQ	0 0:49:3:58,50	0 1:3:5:65,3	0/0:41:3
GT:GQ:DP:HQ	1 2:21:6:23,27	2 1:2:0:18,2	2/2:35:4
GT:GQ:DP:HQ	0 0:54:7:56,60	0 0:48:4:51,51	0/0:61:2
GT:GQ:DP	0/1:35:4	0/2:17:2	1/1:40:3

<https://samtools.github.io/hts-specs/VCFv4.4.pdf>

The VCF format

CHROM	chromosome
POS	position of the start of the variant
ID	unique identifier of the variant (e.g. rs number for SNPs)
REF	reference allele
ALT	comma separated list of alternate non-reference alleles
QUAL	phred-scaled quality score
FILTER	site filtering information
INFO	user extensible annotation (e.g. samtools and GATK may differ in this)

#CHROM	POS	ID	REF	ALT	QUAL	FILTER	INFO	FORMAT	NA00001	NA00002	NA00003
20	14370	rs6054257	G	A	29	PASS	NS=3;DP=14;AF=0.5;DB;H2	GT:GQ:DP:HQ	0 0:48:1:51,51	1 0:48:8:51,51	1/1:43:5:.,.
20	17330	.	T	A	3	q10	NS=3;DP=11;AF=0.017	GT:GQ:DP:HQ	0 0:49:3:58,50	0 1:3:5:65,3	0/0:41:3
20	1110696	rs6040355	A	G,T	67	PASS	NS=2;DP=10;AF=0.333,0.667;AA=T;DB	GT:GQ:DP:HQ	1 2:21:6:23,27	2 1:2:0:18,2	2/2:35:4
20	1230237	.	T	.	47	PASS	NS=3;DP=13;AA=T	GT:GQ:DP:HQ	0 0:54:7:56,60	0 0:48:4:51,51	0/0:61:2
20	1234567	microsat1	GTC	G,GTCT	50	PASS	NS=3;DP=9;AA=G	GT:GQ:DP	0/1:35:4	0/2:17:2	1/1:40:3

<https://samtools.github.io/hts-specs/VCFv4.4.pdf>

Variant Quality vs Genotype Quality

- Variant Quality
 - Phred score estimating if variant is likely to be an artifact
- Genotype Quality
 - Phred score estimating accuracy of estimated sample genotype

#CHROM	POS	ID	REF	ALT	QUAL	FILTER	INFO	FORMAT	NA00001	NA00002	NA00003
20	14370	rs6054257	G	A	29	PASS	NS=3;DP=14;AF=0.5;DB;H2	GT:GQ:DP:HQ	0 0:48:1:51,51	1 0:48:8:51,51	1/1:43:5:.,.
20	17330	.	T	A	3	q10	NS=3;DP=11;AF=0.017	GT:GQ:DP:HQ	0 0:49:3:58,50	0 1:3:5:65,3	0/0:41:3
20	1110696	rs6040355	A	G,T	67	PASS	NS=2;DP=10;AF=0.333,0.667;AA=T;DB	GT:GQ:DP:HQ	1 2:21:6:23,27	2 1:2:0:18,2	2/2:35:4
20	1230237	.	T	.	47	PASS	NS=3;DP=13;AA=T	GT:GQ:DP:HQ	0 0:54:7:56,60	0 0:48:4:51,51	0/0:61:2
20	1234567	microsat1	GTC	G,GTCT	50	PASS	NS=3;DP=9;AA=G	GT:GQ:DP	0/1:35:4	0/2:17:2	1/1:40:3

Variant Calling



- 2 Major types of Variants
 - Single Nucleotide Variants (SNV) and small Indels
 - Smaller than the size of one read
 - Large Structural Variants
 - Usually larger than the size of one read



Single Nucleotide Variants (SNV) and small Indels

Variants detected within reads (smaller than size of read)

- SNVs:
 - Change of a single nucleotide
- Indels:
 - “Small” deletion or amplification





Single Nucleotide Variants (SNV) and small Indels

- Main factors affecting detection of SNVs and Indels
 - Number of reads (coverage supporting a variant)
 - Base quality (affects confidence in the SNVs)
 - PCR amplification bias (can generate duplicates and other biases)
 - Repetitive areas (mostly affects indels, but also affects SNVs)

Duplicated Reads

- Duplicate reads (same fragment) can appear
 - In library preparation during amplification (eg. WES)
 - In the amplification process while sequencing (optical duplicates)

```

AAGCGATG          AGTTGTGT
AAGCGATG TCCATGC  AGTTGTGT
          ACTCCAT   GTTGT
                TGCTCAGTT
                        GTGTGTTT-CA

AAGCGATGACTGCATGCACAGTTGTGTGTTTTCACGTGAC
  geneA          geneB          geneC
    
```

Duplicated Reads: optical duplicates

@MN00723:33:000H3MCVT:1:11102:7591:1087 1:N:0:47

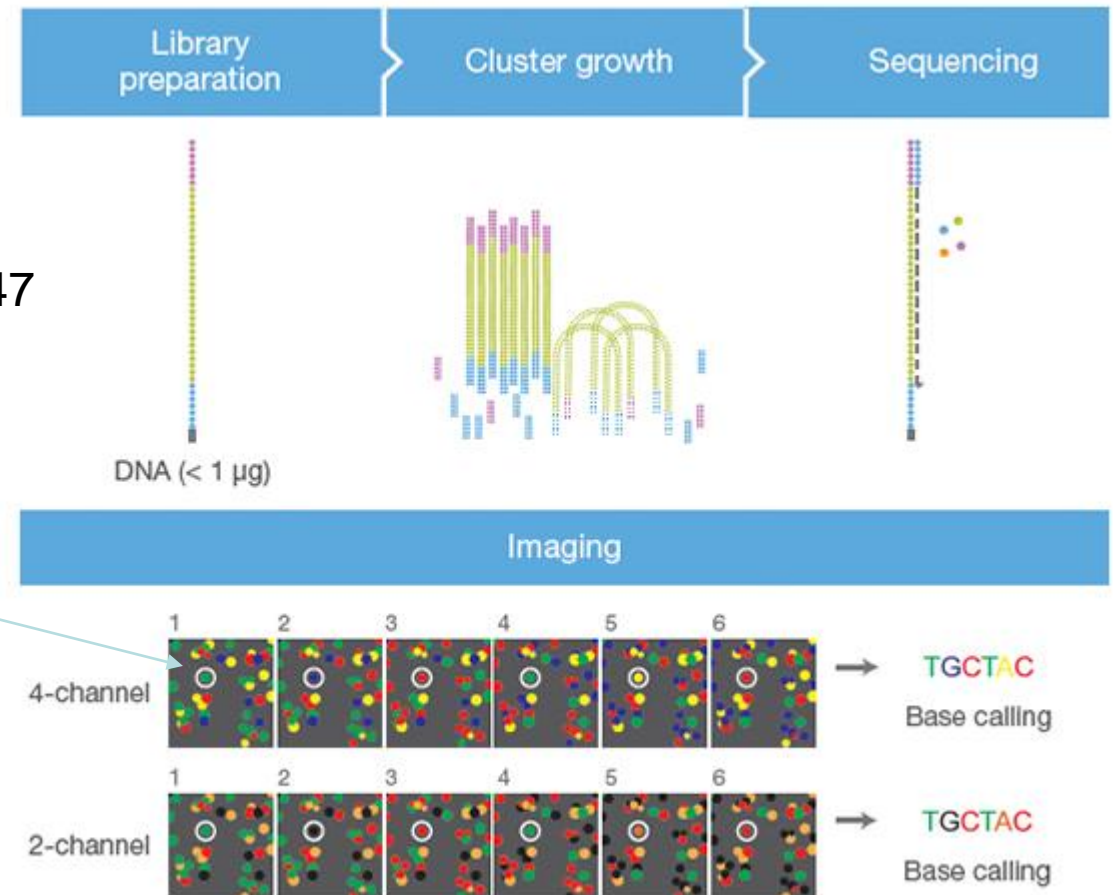
Machine

Flow Cell

Lane

Position

Optical duplicates are duplicate reads that are very close in the flow cell




Duplicated Reads

- **The recommended practice is to ignore duplicates**
 - Only consider one of the duplicates for variant calling
 - Usually the one with the best quality
 - This may remove good information (eg. with high coverage, targeted)
 - Duplicates are marked and later ignored (or not)
 - Benefits of marking duplicates not always obvious
 - <https://www.ncbi.nlm.nih.gov/pmc/articles/PMC4965708/>
 - Eg. one can choose to only ignore reads marked as optical duplicates

Base Quality Recalibration

Base Quality Depends on several factors:

- Sample Quality (DNA)
- Nucleotide context
- Machine and cycle of sequencing
- Type of variant (SNP or Indel)



The machine may not estimate well the base quality score

Base Quality Recalibration

Use list of known variants to estimate correct quality values

- All bases different from reference **not in the provided list** of known variants are considered to be errors




Base Quality Recalibration

Base Quality Recalibration:

The covariates being used here:

- ReadGroupCovariate
- QualityScoreCovariate
- ContextCovariate
- CycleCovariate



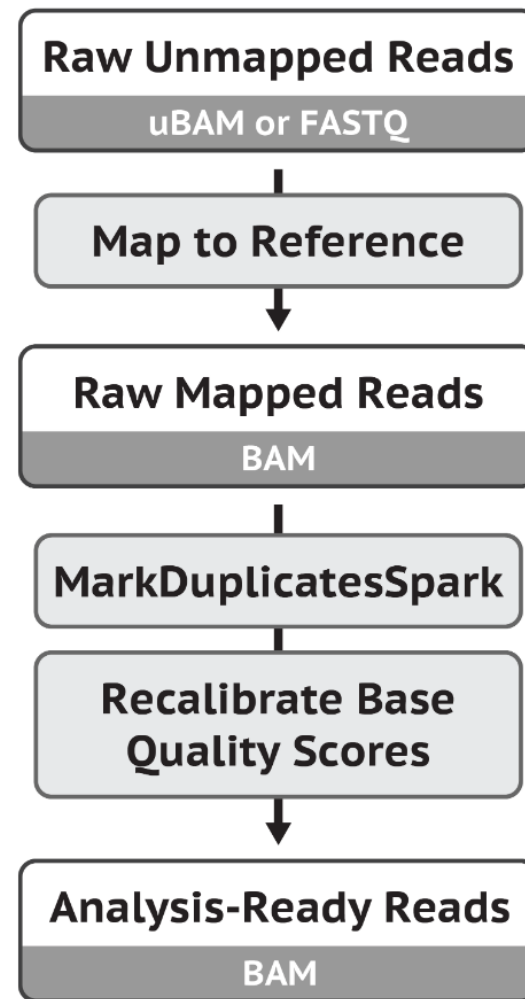
<https://www.ncbi.nlm.nih.gov/pmc/articles/PMC3083463/>

<https://www.youtube.com/watch?v=L4D1dwES9s8>

<https://gatk.broadinstitute.org/hc/en-us/articles/360035890531-Base-Quality-Score-Recalibration-BQSR->

Data pre-processing for variant discovery

GATK
Best
Practices



<https://gatk.broadinstitute.org/hc/en-us/articles/360035535912-Data-pre-processing-for-variant-discovery>

Example estimating Genotype Quality

- These calculations are software-dependent
 - Example Genotype Likelihood (GATK)

$$\mathcal{L}(g) = \frac{1}{m^k} \prod_{j=1}^l [(m-g)\epsilon_j + g(1-\epsilon_j)] \prod_{j=l+1}^k [(m-g)(1-\epsilon_j) + g\epsilon_j]$$

g: genotype (i.e. 0, 1 or 2)

m: ploidy (2 for human)

ϵ : base error

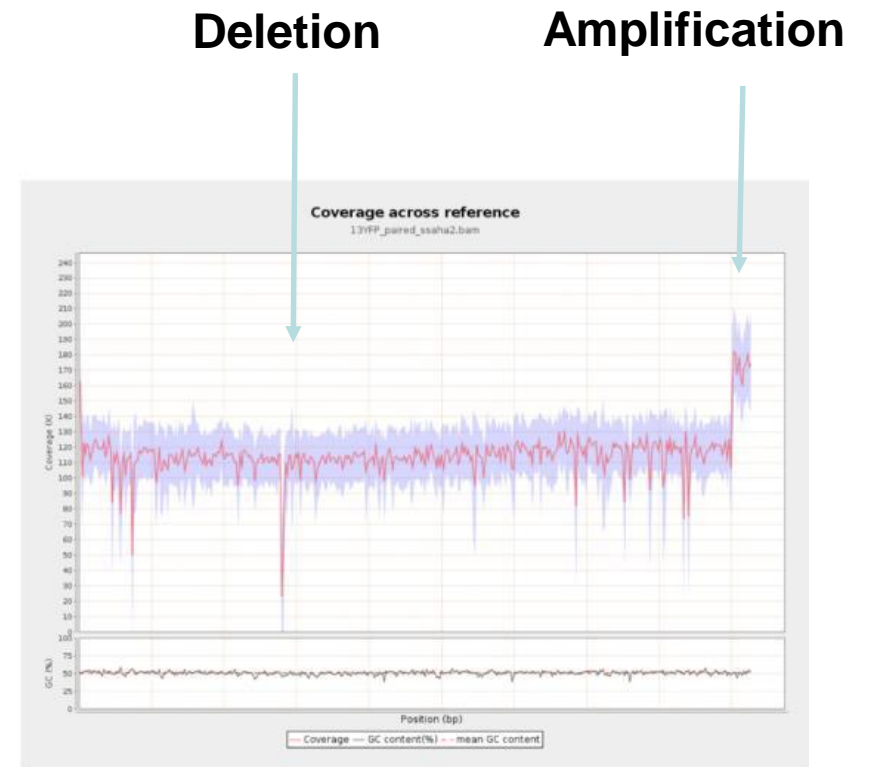
k: number of bases at the site

l: number of bases that equal reference

Large Structural Variants


Variants larger than the size of reads

- Large Deletions and Amplifications
 - Gene Deletions and Duplications
- Other Structural Variants
 - Fusions; Inversions; Transposons...
- Horizontal Transfer
 - Novel genomic regions



Large Structural Variants

- Evidence used to detect Structural Variants
 - Differences in Coverage
 - Most commonly used
 - Particularly with targeted sequencing
 - Although there's still amplification bias



Large Structural Variants

- Evidence used to detect Structural Variants
 - Junction evidence (difficult in targeted sequencing)
 - Can use paired read information (namely, expected fragment length –noisier)
 - Can use information within reads (more precise - requires bigger reads)


```

GCCACGTTACAATACTCTTATCAGTAATAGGCTGGCAAAAACTTT-----
-----GGTAATGACTCCAACCTATTGATAGTGTTTATGTTAGATAATGCCCGA
CCACGTTACAATACTCTTATCAGTAATAGGCTGGCAAAAACTTTGGTAATGACTCCAACCTATTGATAGTGTTTT
CCACGTTACAATACTCTTATCAGTAATAGGCTGGCAAAAACTTTGGTAATGACTCCAACCTATTGATAGTGTTTT
CCACGTTACAATACTCTTATCAGTAATAGGCTGGCAAAAACTTTGGTAATGACTCCAACCTATTGATAGTGTTTTA
CCACGTTACAATACTCTTATCAGTAATAGGCTGGCAAAAACTTTGGTAATGACTCCAACCTATTGATAGTGTTTTAT
CCACGTTACAATACTCTTATCAGTAATAGGCTGGCAAAAACTTTGGTAATGACTCCAACCTATTGATAGTGTTTTATGTT
CCACGTTACAATACTCTTATCAGTAATAGGCTGGCAAAAACTTTGGTAATGACTCCAACCTATTGATAGTGTTTTATGTT

```

Examples of variants


Example mutations



Example deletion




- IGV provides colors to signal unusual situations
 - Besides mutations, information from paired-end is also there



Pairs on different chromosomes

Insert Size Lengths

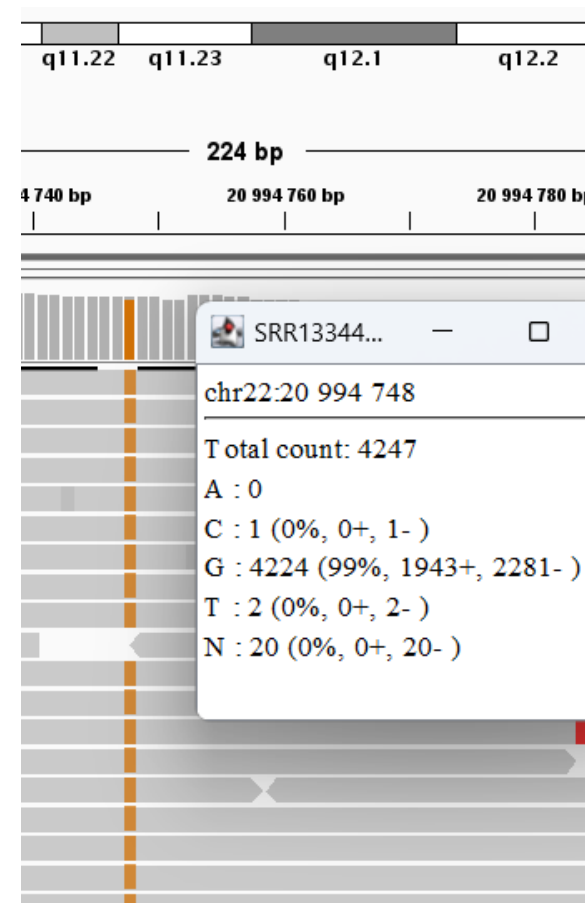
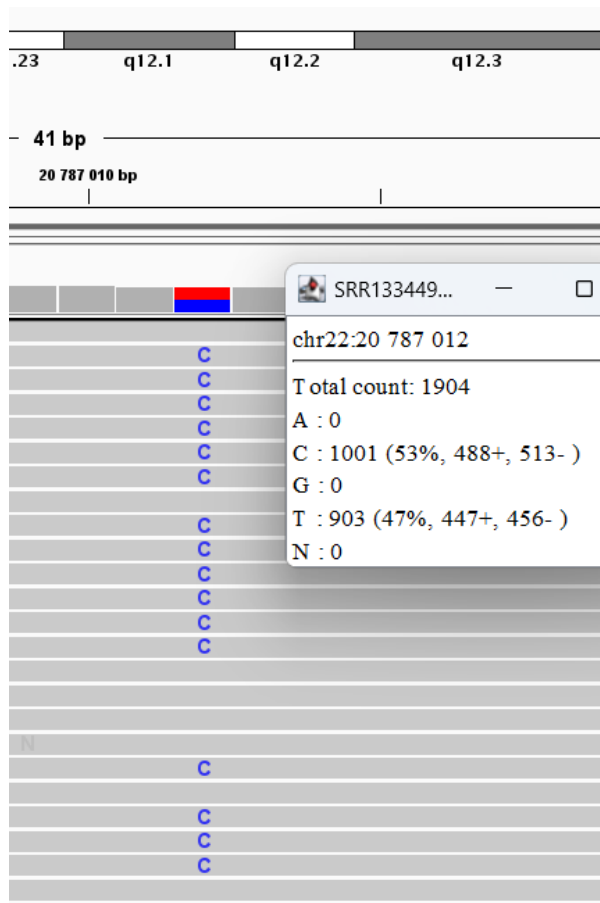


Pair orientation

<https://software.broadinstitute.org/software/igv/>

Visualization of Read Mappings

- Example of mutations



Variant Selection

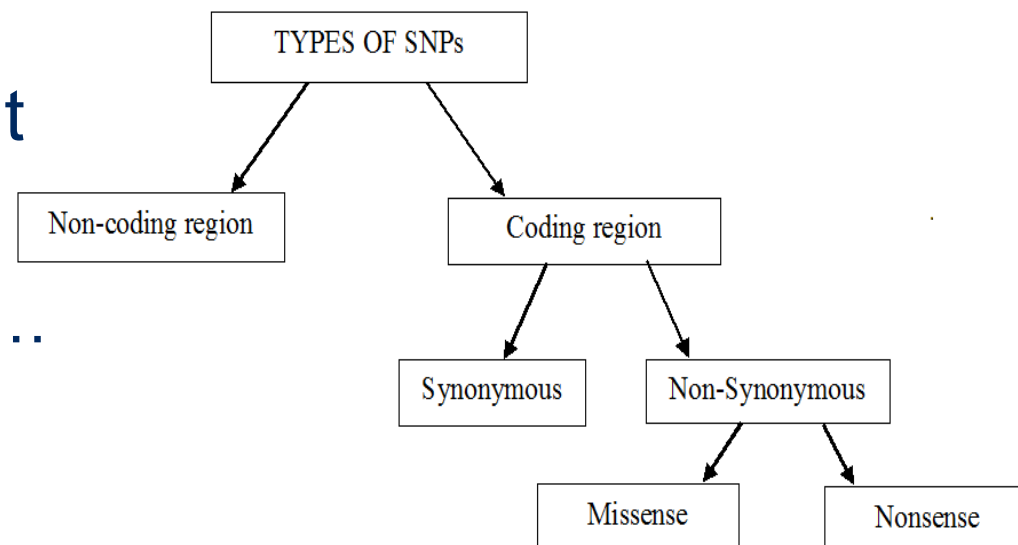
Criteria to select “valid” variants:

- Higher Number of reads (coverage supporting a variant)
- Low Bias in the Base quality supporting the variant
- Low bias in the strand of the reads supporting the variant
- Avoid variants only at the end of reads (repetitive areas)
- Avoid duplicate reads
- Etc...

Variant Annotation

Criteria to select “relevant” variants:

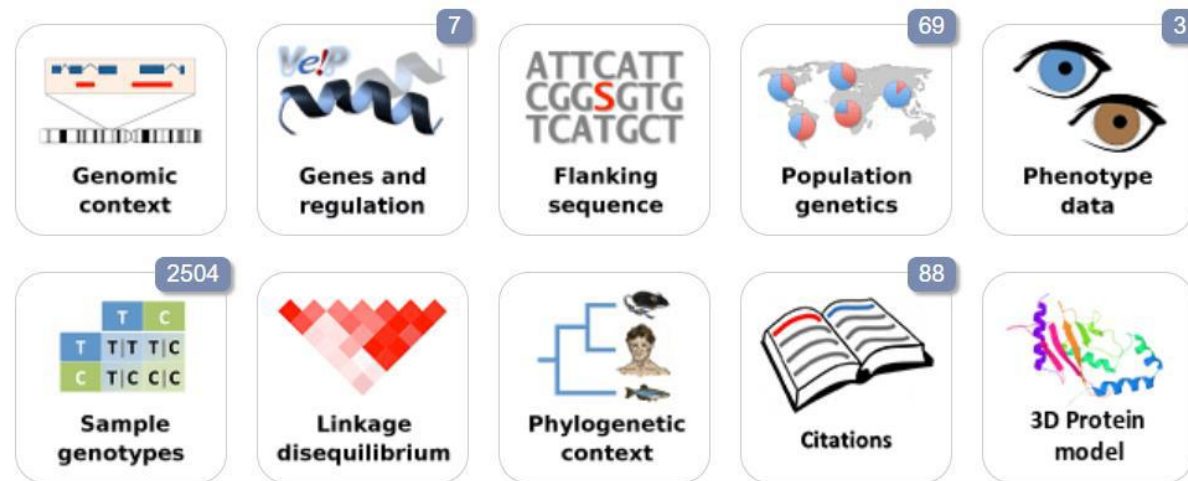
- Coding versus non-coding
 - Coding: Silent versus non-silent
 - Non-coding: can be complex
 - splice-sites, regulatory regions,...



Variant Annotation

Criteria to select “relevant” variants:



- Population frequency
- Disease-association



Only works for species with a lot of data, such as human

Variant Annotation

- Criteria to select “relevant” variants:
 - Simple exemple with SARS-CoV-2 (Omicron)



A few optional references:

- Single Nucleotide Variants and small Indels
 - <https://www.ncbi.nlm.nih.gov/pmc/articles/PMC3083463/>
- DeepVariant: Using AI to find variants
 - <https://github.com/google/deepvariant>; <https://www.nature.com/articles/nbt.4235>
- Copy Number Alterations and other Structural Variants
 - <https://www.ncbi.nlm.nih.gov/pmc/articles/PMC4300727/>
- Finding clonal vs subclonal variants
 - <https://www.sciencedirect.com/science/article/pii/S2001037017300946>
- Variant annotation
 - <https://www.nature.com/articles/nprot.2015.105>
 - <https://genomebiology.biomedcentral.com/articles/10.1186/s13059-016-0974-4>



Practical Exercise

<https://github.com/dsobral/MBCB/blob/main/README.md>